
Introduction to Hive
Lecture 27

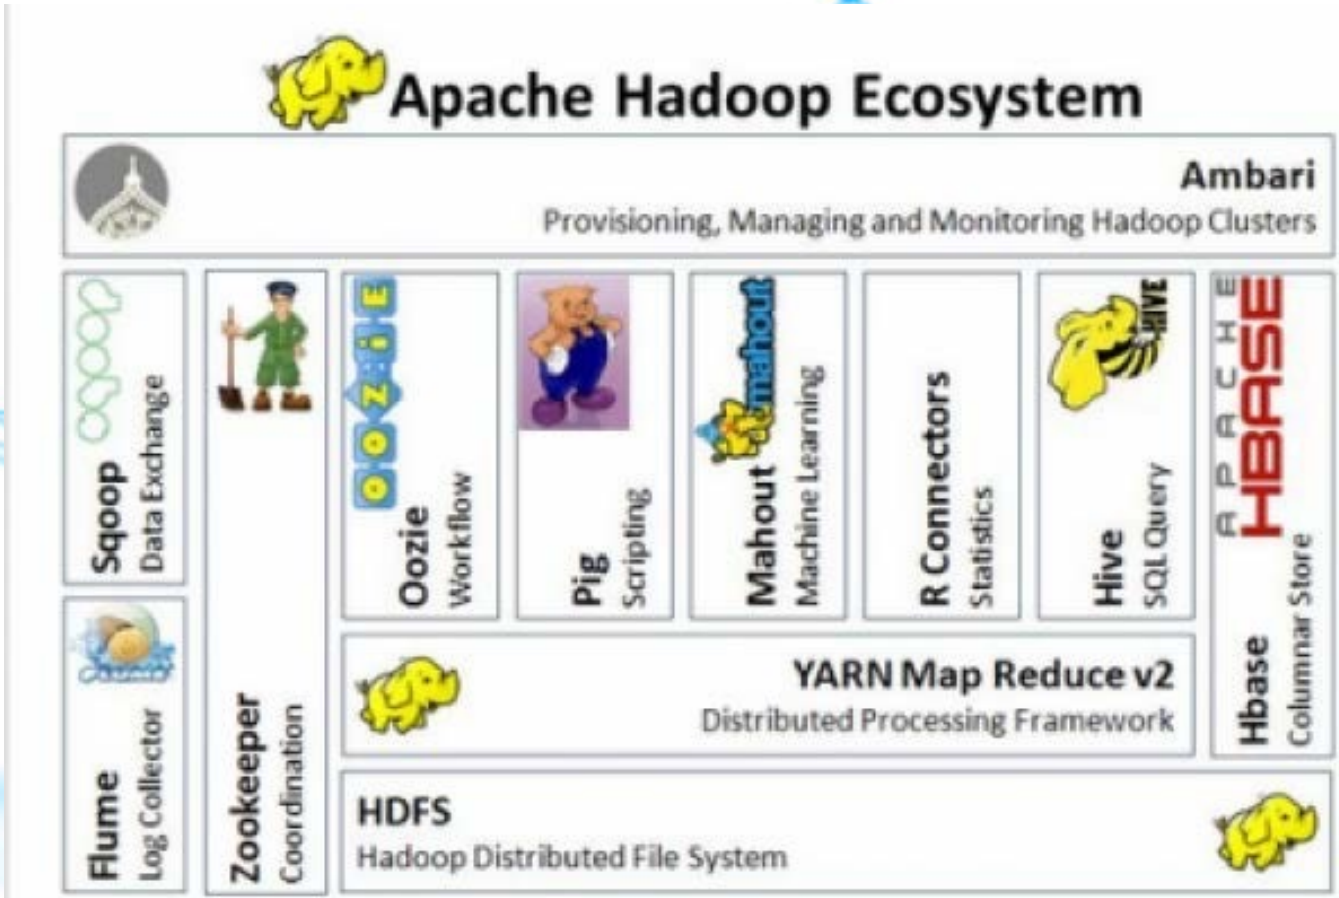
Apache Hadoop Ecosystem
Ambari
Provisioning, Managing and Monitoring Hadoop Clusters
Sqoop
Data Exchange
Flume
Log Collector
Zookeeper
Coordination
Oozie
Workflow
Pig
Scripting
Mahout
Machine Learning
R Connectors
Statistics
Hive
SQL Query
Hbase
Columnar Store
YARN Map Reduce v2
Distributed Processing Framework
HDFS
Hadoop Distributed File System

What is Hive?

Data warehouse infrastructure build on top of Hadoop for querying and managing large data sets

Why Hive?

Hadoop is great!

MapReduce is very low level

Lack of expressiveness

Higher level data processing languages are needed

Hive Features

Designed for OLAP

SQL type language for querying

It is familiar, fast, scalable, and extensible

On Line Analytical Processing

HiveQL or HQL

Can plug in map/reduce scripts in language of choice

Hive is NOT

Relational database

Designed for Online Transaction Processing (OLTP)

*Online transaction processing, or **OLTP**, is a class of information systems that facilitate and manage transaction-oriented applications, typically for data entry and retrieval transaction processing (Wikipedia)*

History

ETL = Extract, Transform and Load

Early Hive development work started at Facebook in 2007

Hive is an Apache project under Hadoop

http://hive.apache.org

Data Warehouse Infrastructure for Hadoop

SQL-like query language (QL)

Enables developers to utilize custom mappers and reducers

Provides tools to enable ETL on large data

Hive Architecture and Components

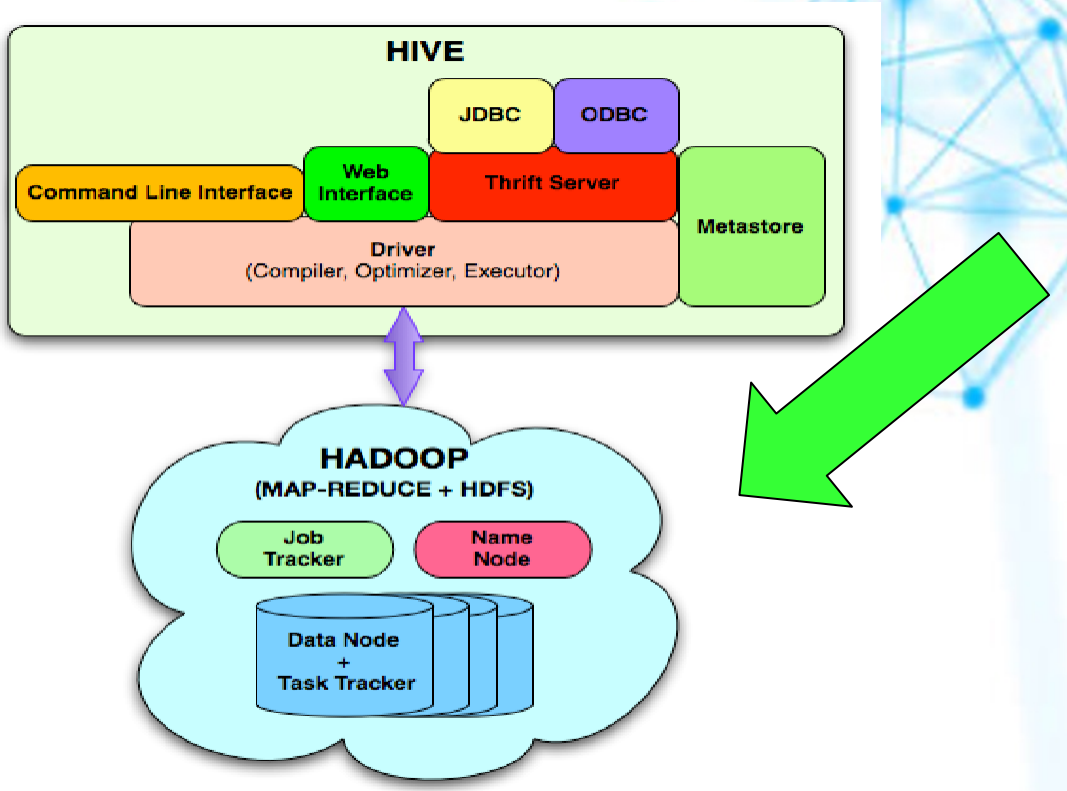

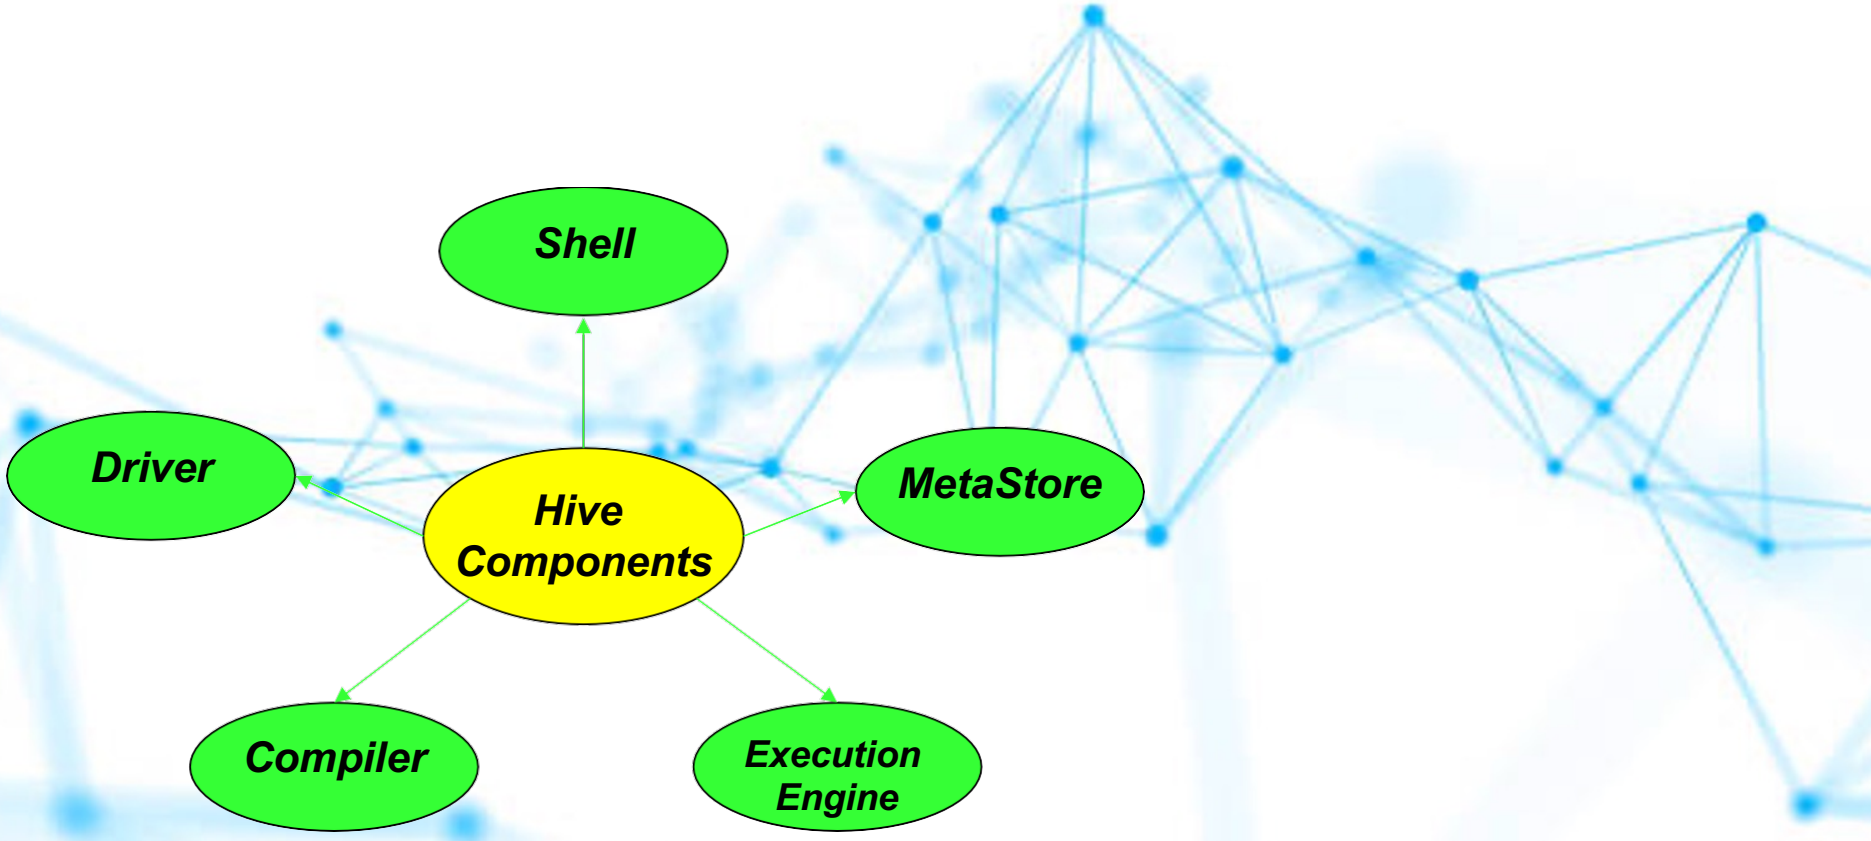
Shell
Driver
Hive Components
MetaStore
Compiler
Execution Engine

Hive Architecture

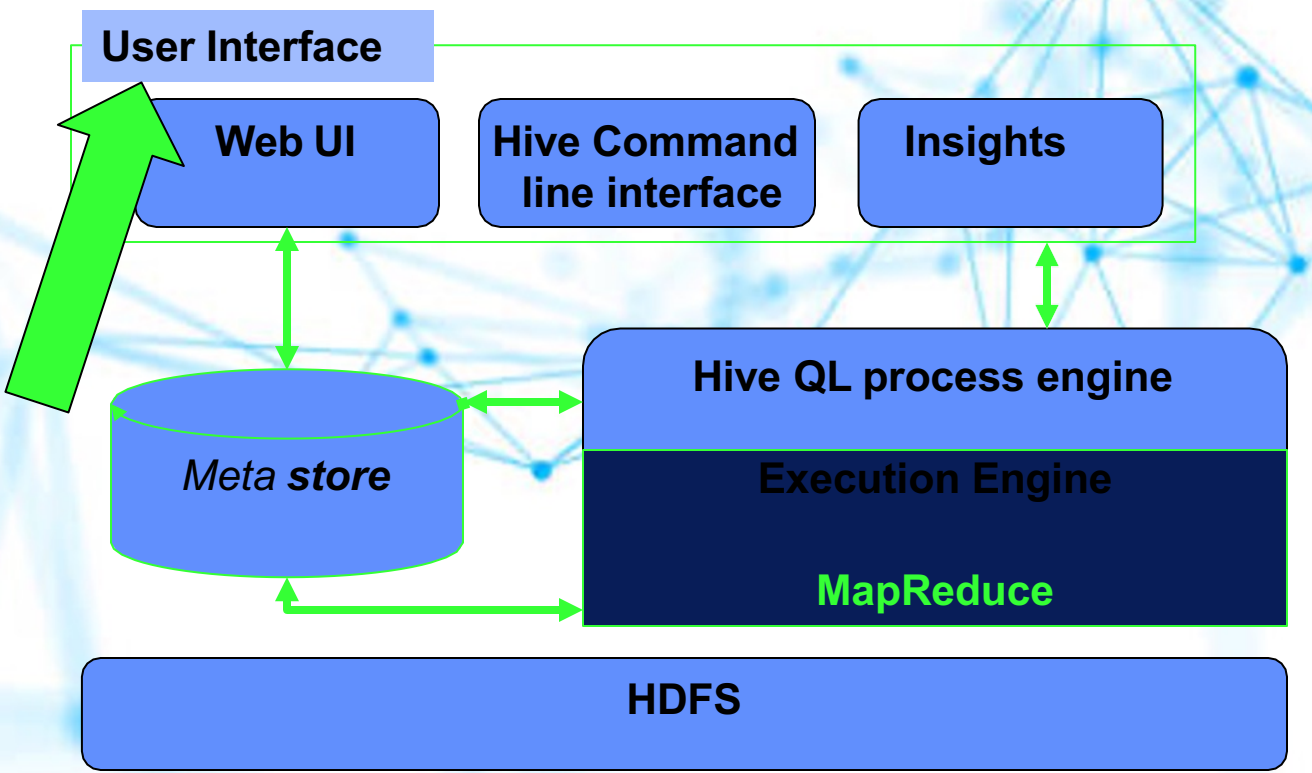

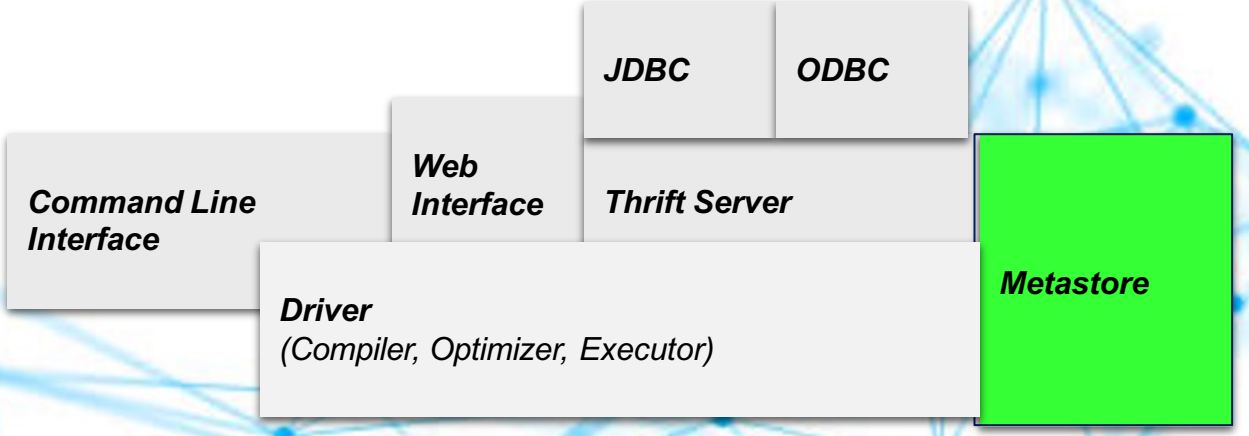

Metastore

Stores the system catalog and meta data about tables, columns, partitions etc.

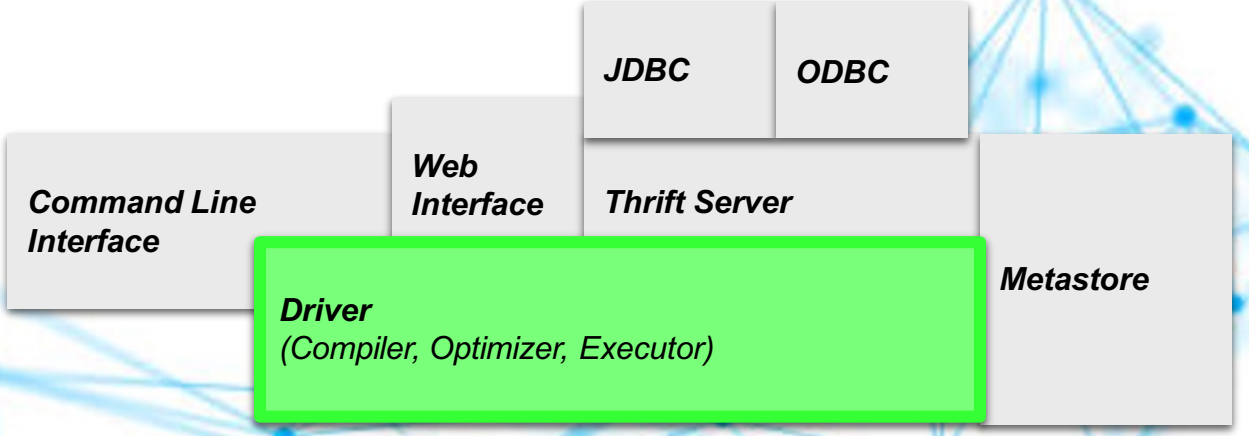

Driver

- Manages the lifecycle of a HiveQL statement
- Maintains a session handle and any session statistics

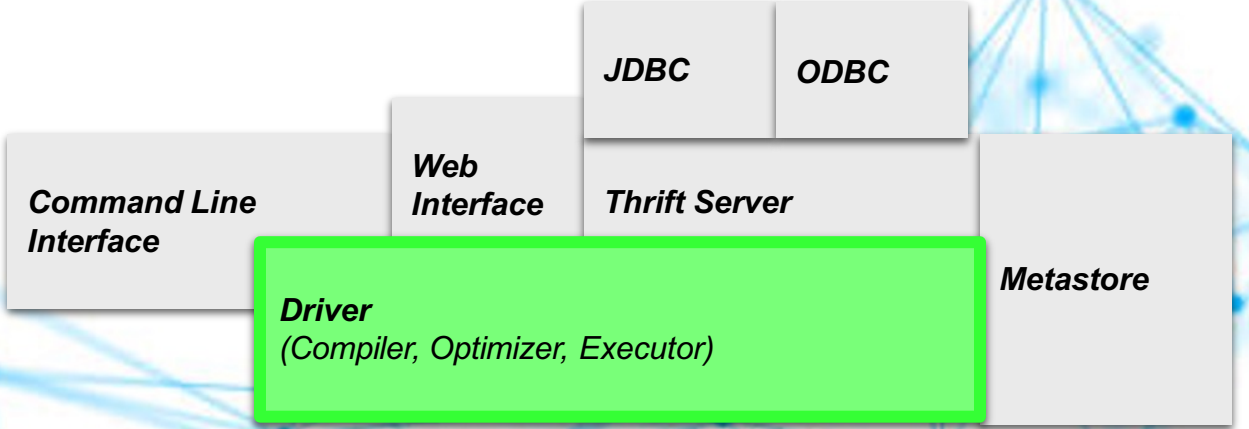

Query Compiler

The component that compiles HiveQL into a directed acyclic graph of map/reduce tasks

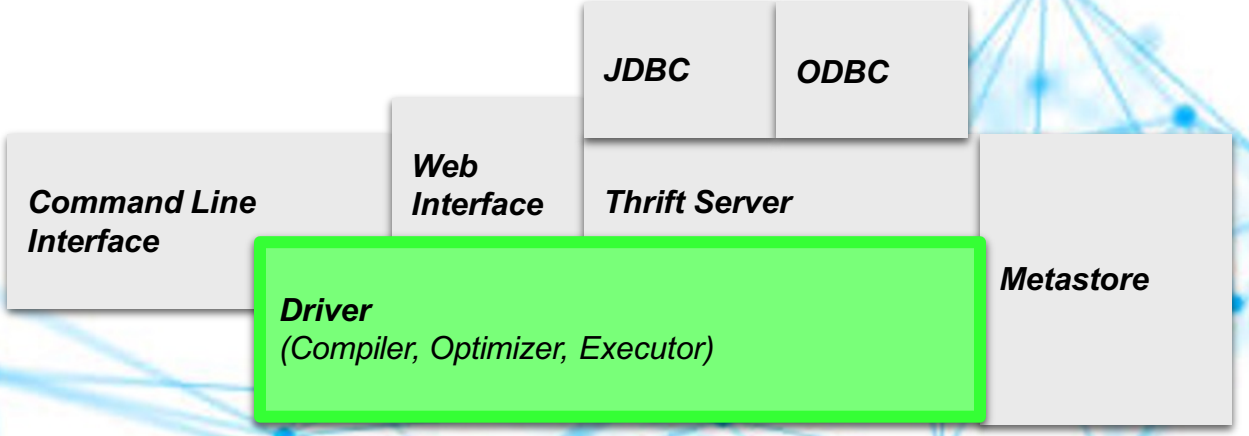

Optimizer

Consists of a chain of transformations

Performs Column Pruning , Partition Pruning, Repartitioning of Data

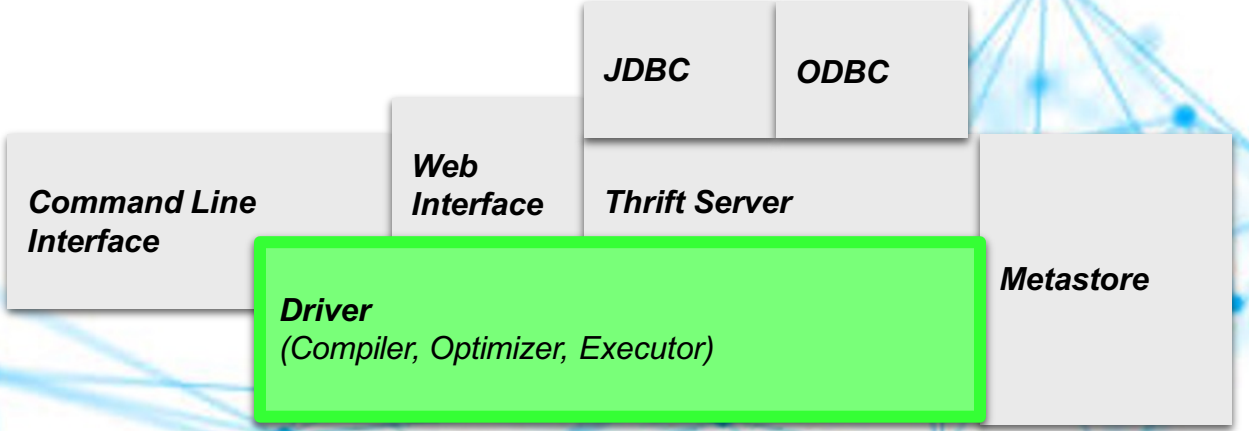

Execution Engine

- Executes the tasks produced by the compiler in proper dependency order
- Interacts with the underlying Hadoop instance

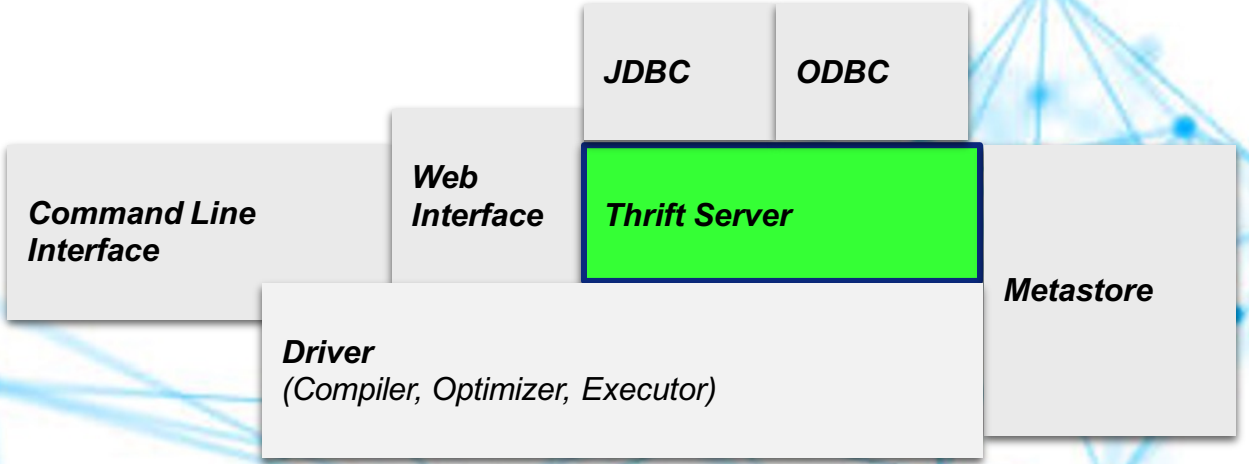

HiveServer

Provides a Thrift interface and a JDBC/ODBC server Enables Hive integration with other applications

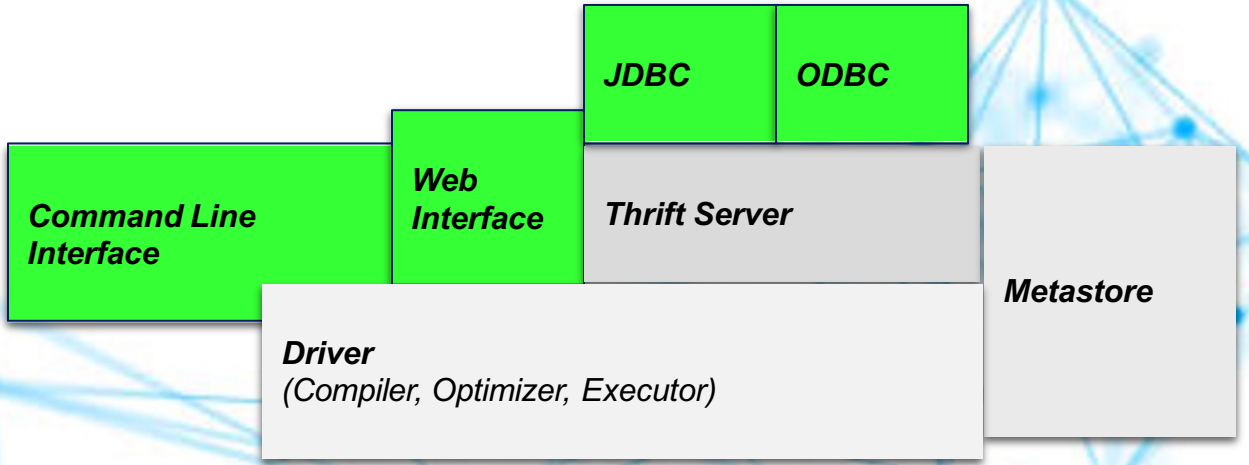

Client Components

- Command Line Interface(CLI)
- Web UI
- JDBC/ODBC driver

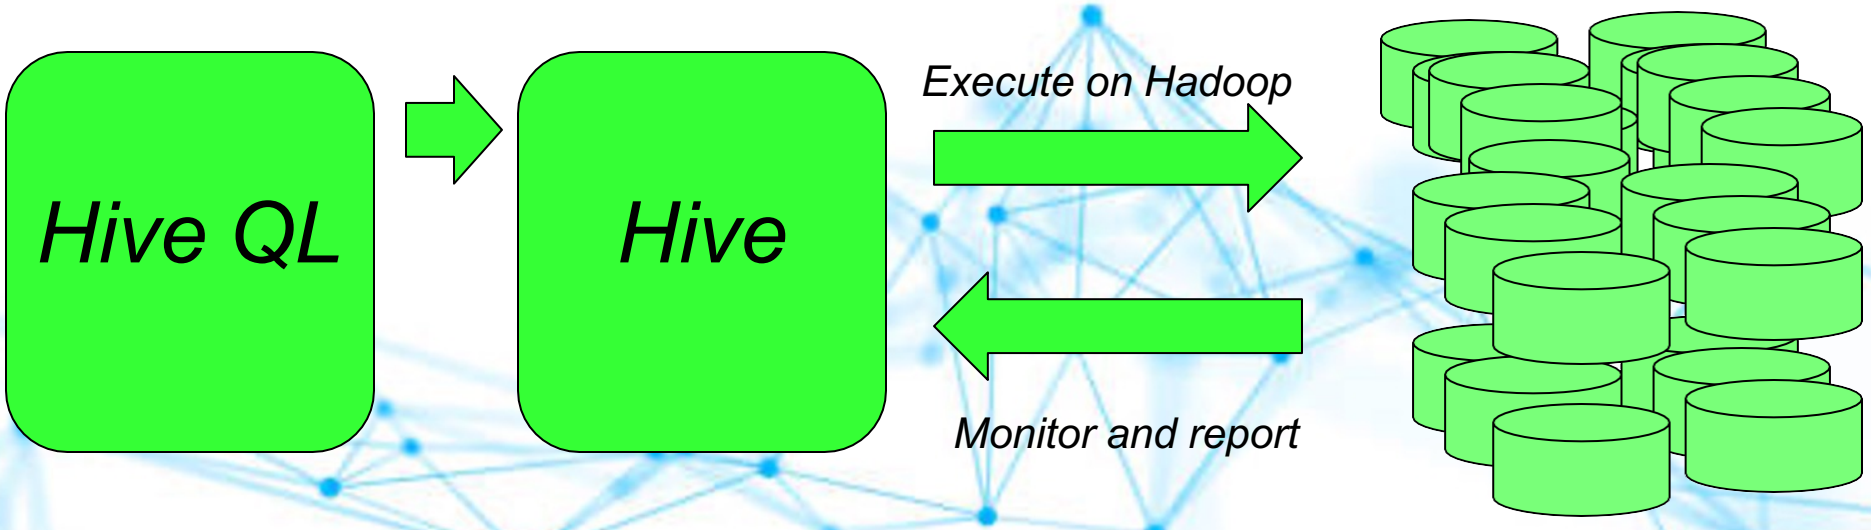
Hive QL
Hive
Execute on Hadoop
Monitor and report

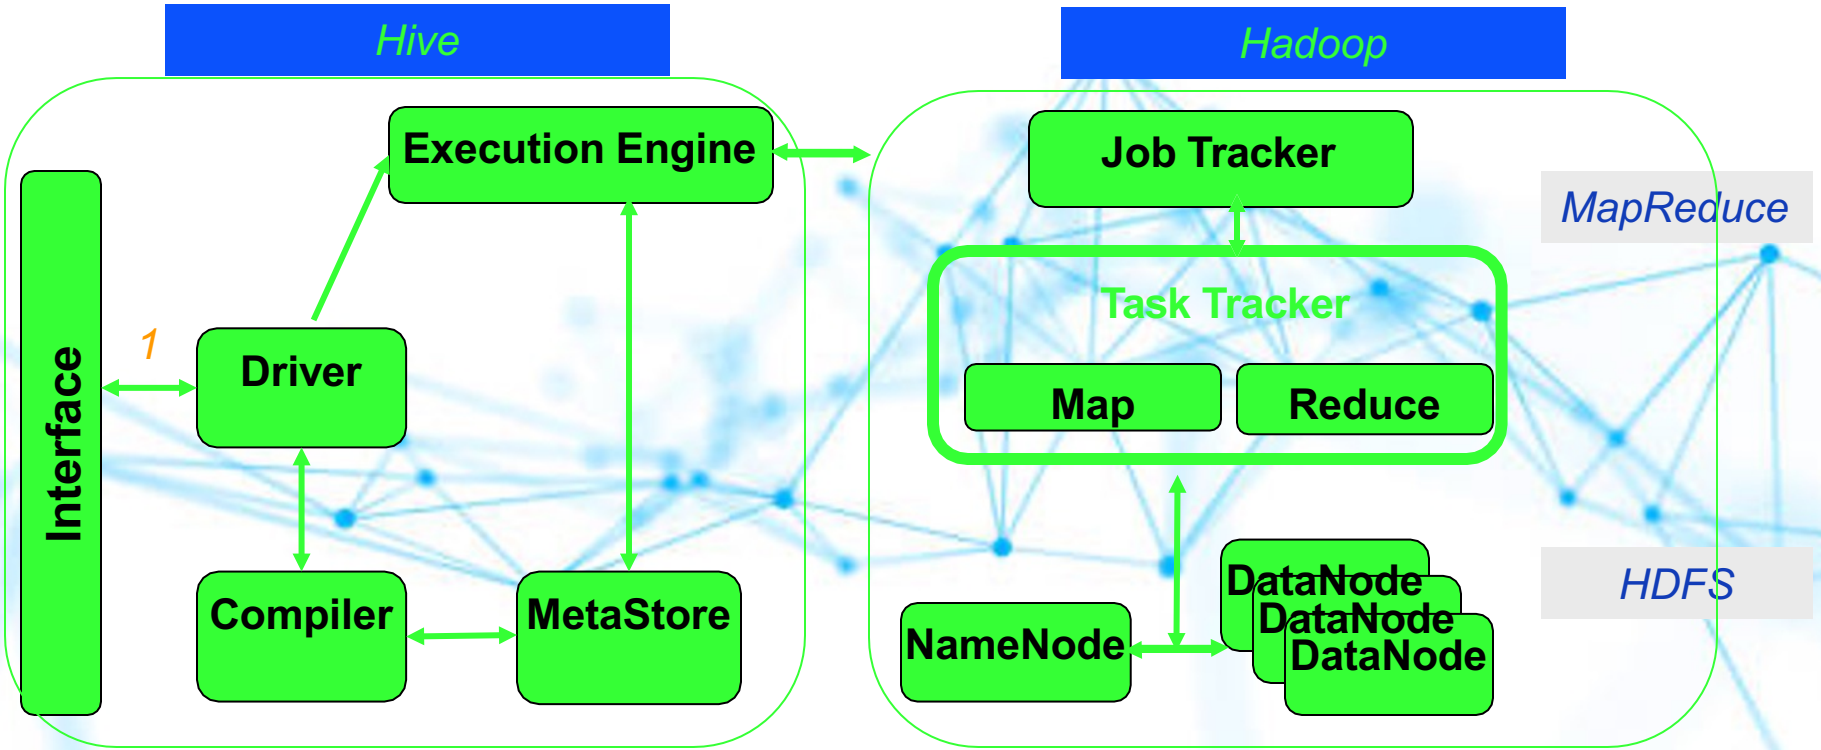
Hive
Hadoop
Execution Engine
Job Tracker
MapReduce
Interface
1
Driver
Task Tracker
Map
Reduce
Compiler
MetaStore
DataNode
DataNode
DataNode
NameNode
HDFS

Hive's Data Units

- Databases
- Tables
- Partitions
- Buckets (or clusters)

Very similar to SQL and Relational DBs

3-Levels: Tables → Partitions → Buckets

Data Model

Table maps to a HDFS directory

Partition maps to sub-directories under the table

Bucket maps to files under each partition

Tables

Similar to tables in relational DBs

Each table has corresponding directory in HDFS

Partitions

- Analogous to dense indexes on partition columns
- Nested sub-directories in HDFS for each combination of partition column values
- Allows users to efficiently retrieve rows

Hive Data Structures

Traditional Database concepts

Supports primitive types

Additional types and structures

Hive Data Structures

Traditional database concepts

- Tables
- Rows
- Columns
- Partitions

Hive Data Structures

Basic types

- Integers
- Floats
- Doubles
- Strings

Hive File Formats

Hive enables users store different file formats

Performance improvements

- TEXTFILE
- SEQUENCEFILE
- ORC
- RCFILE

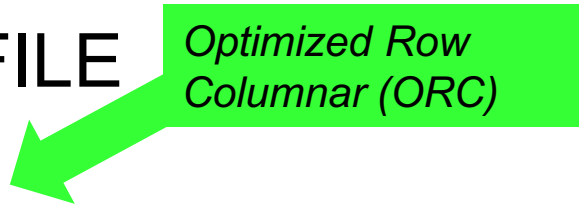

Hive Commands

Hive Interface

Command Line interface

Web interface or Hue

Java Database connectivity

Hive Commands

Database

Set of Tables - name conflicts resolution

Table

Set of Rows - have the same columns

Row

A single record - a set of columns

Column

Value and type for a single value

Tables Commands

- SHOW TABLES
- CREATE TABLE
- ALTER TABLE
- DROP TABLE

Hive Commands

CREATE TABLE mytable (myint INT, bar STRING)
PARTITIONED BY (ds STRING);

SHOW TABLES '.*my';

ALTER TABLE mytable **ADD COLUMNS** (new_col INT);

DROP TABLE mytable;

A table in Hive is an HDFS directory in Hadoop

Hive Commands

Schema is known at creation time (like DB schema)

Partitioned tables have “sub-directories”, one for each partition

```
CREATE TABLE mypeople (
id              int,
name            string
)
partitioned by (date string)
```