



SQL: SCHEMA DEFINITION, BASIC CONSTRAINTS, AND QUERIES

SQL INTRODUCTION

- SQL stands for **Structured Query Language**
- Standard language for querying and manipulating data
- Data Definition Language (DDL)
 - Create/alter/delete tables and their attributes
- Data Manipulation Language (DML)
 - Query one or more tables
 - Insert/delete/modify tuples in tables

Many standards out there: ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3), ...

Table name

TABLES IN SQL

Department

Attribute names

DName	Dnumber	MgrSsn	Mgrstartdate
Gizmo	19	324521	1992-08-11
Powergizmo	29	624545	1982-01-21
SingleTouch	1	986133	1912-04-02
MultiTouch	2	1123455	2002-08-28

Tuples or rows

CREATE TABLE

- Creates a new relation(table) in the database
 - Specifies relation's attributes and their data types
 - Specifies constraints such as NOT NULL , UNIQUE ,CHECK etc...

```
CREATE TABLE DEPARTMENT
```

```
(DNAME          VARCHAR(10)    NOT NULL ,  
 DNUMBER        INTEGER CHECK(DNUMBER >0 AND DNUMBER <25),  
 MGRSSN         CHAR(9),  
 MGRSTARTDATE   DATE  
);
```

CREATE SCHEMA

- Specifies a new database schema by giving it a name
- Example:

CREATE SCHEMA COMPANY AUTHORIZATION Zareen;

Data Types

○ Numeric

- integer
- Smallint
- Floating point
- Decimal(Precision, Scale)
- Numeric(Precision,Scale)

○ Character String

- Char(n)
- Varchar(n)
- Nvarchar(n)

○ Bit-String

○ Boolean



ADDITIONAL DATA TYPES

○ **DATE:**

- Made up of year-month-day in the format yyyy-mm-dd

○ **TIME:**

- Made up of hour:minute:second in the format hh:mm:ss

○ **TIME(i):**

- Made up of hour:minute:second plus i additional digits specifying fractions of a second
- format is hh:mm:ss:ii...i

○ **TIMESTAMP:**

- Has both DATE and TIME components

CREATE DOMAIN

- We can declare a domain in SQL
- This makes it easier to change the data type of numerous attributes in a schema
- Improves Schema readability
- Not available in many SQL implementation like (T-SQL)

CREATE DOMAIN SSN_TYPE AS CHAR(9);

CONSTRAINTS IN SQL

- CREATE TABLE command allows us to specify the primary key, secondary keys, and foreign keys.
- Key attributes can be specified via the PRIMARY KEY and UNIQUE phrases

```
CREATE TABLE  DEPARTMENT
(
    DNAME          VARCHAR(10)    NOT NULL,
    DNUMBER        INTEGER        NOT NULL,
    MGRSSN         CHAR(9)        NULL,
    MGRSTARTDATE   CHAR(9),
    PRIMARY KEY (DNUMBER),
    UNIQUE (DNAME),
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE
);
```

REFERENTIAL INTEGRITY OPTIONS

- We can specify RESTRICT, CASCADE, SET NULL or SET DEFAULT on foreign keys.

```
CREATE TABLE DEPARTMENT
(
    DNAME          VARCHAR(10)    NOT NULL,
    DNUMBER        INTEGER        NOT NULL,
    MGRSSN         CHAR(9) ,
    MGRSTARTDATE   CHAR(9),
    PRIMARY KEY (DNUMBER),
    UNIQUE (DNAME),
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE
    ON DELETE SET DEFAULT ON UPDATE CASCADE
);
```

CASCADE : to delete a row with a key referenced by foreign keys in existing rows in other tables, all rows that contain those foreign keys are also deleted.

REFERENTIAL INTEGRITY OPTIONS

Employee

<u>ssn</u>	supervisor
123456789		234589710
... ..		
234589710		null

delete

Employee



<u>ssn</u>	supervisor
123456789		234589710
... ..		
234589710		null

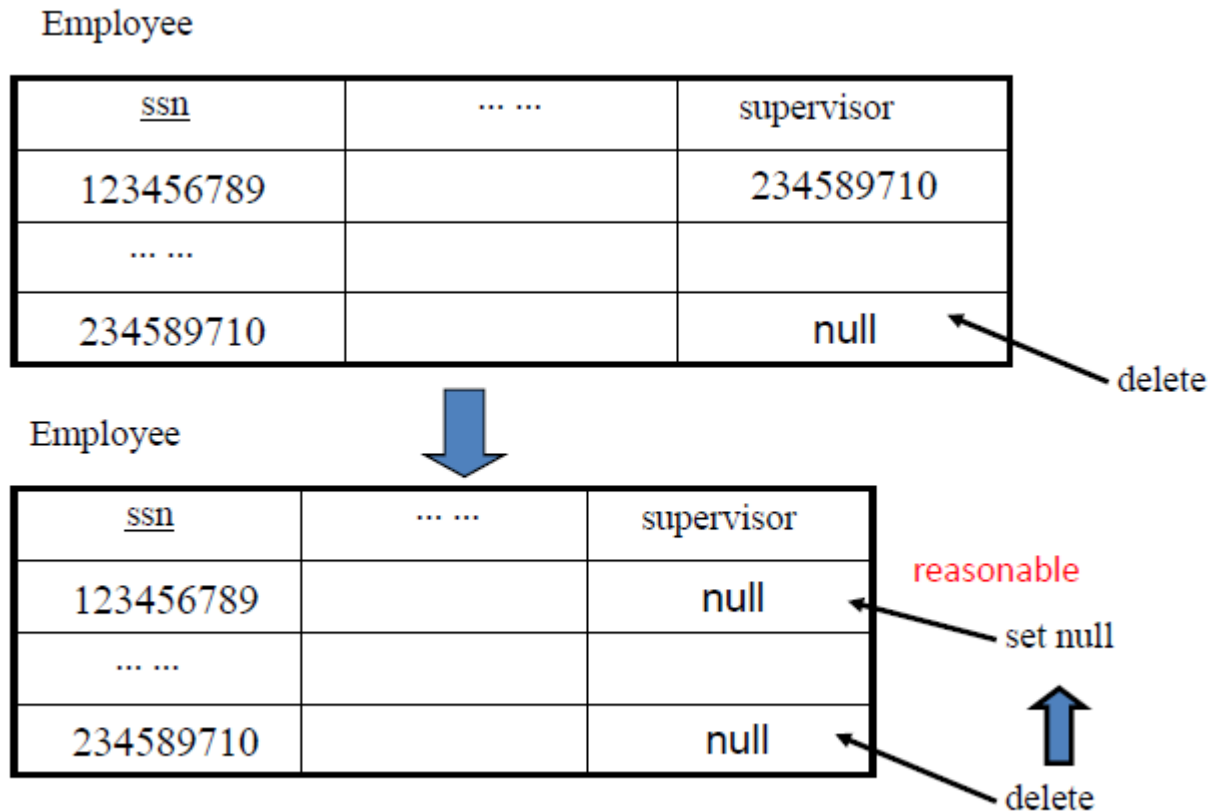
not reasonable

delete

cascade

delete

REFERENTIAL INTEGRITY OPTIONS



REFERENTIAL INTEGRITY OPTIONS

```
CREATE TABLE EMPLOYEE
(
    ENAME VARCHAR(30)    NOT NULL,
    ESSN   CHAR(9),
    BDATE  DATE,
    DNO    INTEGER DEFAULT 1,
    SUPERSSN CHAR(9),
    PRIMARY KEY (ESSN),
    FOREIGN KEY (DNO) REFERENCES DEPARTMENT
    ON DELETE SET DEFAULT ON UPDATE CASCADE,
    FOREIGN KEY (SUPERSSN) REFERENCES EMPLOYEE
    ON DELETE SET NULL ON UPDATE CASCADE
);
```

SQL CONSTRAINTS

- Assigning Names to Constraints

```
CONSTRAINT deptPK PRIMARY KEY(Dnumber)  
CONSTRAINT deptSK UNIQUE(Dname)
```

- CHECK Constraint

```
CHECK (Dept_create_date <= Mgr_start_date)
```

DROP COMMAND

- Drop Command is used to delete schema or named schema elements such as table, domains, or constraints
- Example:

DROP TABLE DEPENDENT;

DROP TABLE EMPLOYEE CASCADE;

DROP SCHEMA COMPANY;

In SQL-Server (T-SQL), DROP TABLE cannot be used to drop a table that is referenced by a FOREIGN KEY. The referencing FOREIGN KEY or the referencing table must first be dropped.

ALTER COMMAND

- The definition of table or named schema elements can be changed using ALTER command
- ALTER can be used to add an attribute to the relation
 - Initially, the new attribute will have NULLs in all the tuples of the relation
 - NOT NULL constraint is *not allowed* for such an attribute
- **Example :**
ALTER TABLE EMPLOYEE ADD COLUMN JOB VARCHAR(12);
T-SQL syntax
ALTER TABLE EMPLOYEE ADD JOB VARCHAR(12);
 - The database user have to enter a value for the new attribute JOB for each EMPLOYEE tuple.

ALTER TABLE

- ALTER command can be use to add or drop constraints

- **Example :**

ALTER TABLE EMPLOYEE add constraint unEmp UNIQUE(NAME) ;

ALTER TABLE EMPLOYEE drop constraint unEmp ;

SQL QUERIES

Basic form:

```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```


- *Not same as* the SELECT operation of the relational algebra
- The result can have duplicate tuples
- SQL relation is a *multi-set* (bag) of tuples; *not* a set of tuples

SIMPLE SQL QUERY

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE category='Gadgets'
```



PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks


“selection”

SIMPLE SQL QUERY

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Price > 100
```




“selection” and
“projection”

PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi



ELIMINATING DUPLICATES


```
SELECT DISTINCT category  
FROM Product
```



Category
Gadgets
Photography
Household

Compare to:

```
SELECT category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household



RELATIONAL DATABASE SCHEMA

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

SIMPLE SQL QUERIES

- Basic SQL queries correspond to using the SELECT, PROJECT, and JOIN operations of the relational algebra
- Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

```
SELECT  BDATE, ADDRESS  
FROM    EMPLOYEE  
WHERE   FNAME='John' AND MINIT='B'  
        AND LNAME='Smith'
```

- Similar to a SELECT-PROJECT pair of relational algebra operations

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

JOIN OPERATION

- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT  FNAME, LNAME, ADDRESS
FROM    EMPLOYEE, DEPARTMENT
WHERE   DNAME='Research' AND DNUMBER=DNO
```

- DNAME='Research' is a *selection condition*
- DNUMBER=DNO is a *join condition*

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

JOIN(CONT.)

- For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```

SELECT PNUMBER, DNUM, LNAME, BDATE, ADDRESS
FROM   PROJECT, DEPARTMENT, EMPLOYEE
WHERE  DNUM=DNUMBER AND MGRSSN=SSN AND
       PLOCATION='Stafford'
    
```

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

UNSPECIFIED WHERE-CLAUSE

- *Missing WHERE-clause*
 - indicates there is no condition and is same as WHERE TRUE
- Retrieve the SSN values for all employees.

```
SELECT SSN
FROM EMPLOYEE
```

- If there is no join condition, then we get *CARTESIAN PRODUCT*

```
SELECT SSN, DNAME
FROM EMPLOYEE, DEPARTMENT
```

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

USE OF *

- To retrieve all the attribute values of the selected tuples, a * is used, which stands for *all the attributes*

Examples:

```
SELECT *  
FROM EMPLOYEE  
WHERE DNO=5
```

```
SELECT *  
FROM EMPLOYEE, DEPARTMENT  
WHERE DNAME='Research' AND DNO=DNUMBER
```

ALIASES

- In SQL, we can use the same name for two (or more) attributes as long as the attributes are in *different relations*
- A query that refers to two attributes with the same name must *prefix* the relation name to the attribute name

Example:

EMPLOYEE.DNO, DEPARTMENT.DNUMBER

ALIASES

- For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.
 - **SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM EMPLOYEE E
WHERE E.SUPERSSN=S.SSN**
- Can also use the AS keyword to specify aliases
 - **SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM EMPLOYEE AS E, EMPLOYEE AS S
WHERE E.SUPERSSN=S.SSN**

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

ARITHMETIC OPERATIONS

- Arithmetic operators '+', '-', '*', and '/') can be applied to numeric values in an SQL query result
- Give all employees who work on the 'ProductX' project a 10% raise.
 - **SELECT FNAME, LNAME, 1.1*SALARY
FROM EMPLOYEE, WORKS_ON, PROJECT
WHERE SSN=ESSN AND PNO=PNUMBER AND
PNAME='ProductX'**

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

ORDER BY

- The **ORDER BY** clause sort the tuples in a query result
- Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name, then first name.

```
SELECT      DNAME, LNAME, FNAME, PNAME
FROM        DEPARTMENT, EMPLOYEE, WORKS_ON, PROJECT
WHERE       DNUMBER=DNO AND SSN=ESSN AND PNO=PNUMBER
ORDER BY    DNAME, LNAME, FNAME
```

- The default order is in ascending order of values
- We can specify the keyword **DESC** if we want a descending order
 - **ORDER BY** Dname **DESC**, Lname **ASC**

NULLS IN SQL QUERIES

- SQL allows queries that check if a value is NULL
- SQL uses **IS** or **IS NOT** to compare NULLs
- Retrieve the names of all employees who do not have supervisors.

- | | |
|---------------|-------------------------|
| SELECT | FNAME, LNAME |
| FROM | EMPLOYEE |
| WHERE | SUPERSSN IS NULL |

- Note: If a join condition is specified, tuples with NULL values for the join attributes are not included in the result

SUBSTRING COMPARISON

- **LIKE** operator is used to compare partial strings
- Two reserved characters are used:
 - '%' (or '*' in some implementations) replaces an arbitrary number of characters, and
 - '_' replaces a single arbitrary character
- Retrieve all employees whose address is in Houston, Texas.
- | | |
|---------------|------------------------------------|
| SELECT | FNAME, LNAME |
| FROM | EMPLOYEE |
| WHERE | ADDRESS LIKE '%Houston,TX%' |

SUBSTRING COMPARISON (CONT.)

- Retrieve all employees who were born during the 1950s.
- Here, '5' must be the third character of the string , so the BDATE value is ' __5_____',.

```
SELECT      FNAME, LNAME  
FROM        EMPLOYEE  
WHERE       BDATE LIKE  ' __5_____',.
```

- LIKE operator allows us to get around the fact that each value is considered atomic and indivisible;
 - hence, in SQL, character string attribute values are not atomic

JOINED RELATIONS IN SQL

- Allows the user to specify different types of joins (regular "theta" JOIN, NATURAL JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, CROSS JOIN, etc)

- **Example:**

- SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM EMPLOYEE AS E , EMPLOYEE AS S
WHERE E.SUPERSSN=S.SSN
- SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM (EMPLOYEE E **JOIN** EMPLOYEE S
ON E.SUPERSSN=S.SSN
- SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM (EMPLOYEE E **LEFT OUTER JOIN** EMPLOYEE S
ON E.SUPERSSN=S.SSN)

JOINED RELATIONS FEATURE IN SQL

```
SELECT  FNAME, LNAME, ADDRESS  
FROM    EMPLOYEE, DEPARTMENT  
WHERE   DNAME='Research' AND DNUMBER=DNO
```

could be written as:

```
SELECT  FNAME, LNAME, ADDRESS  
FROM    (EMPLOYEE JOIN DEPARTMENT  
         ON DNUMBER=DNO)  
WHERE   DNAME='Research'
```

or as:

```
SELECT  FNAME, LNAME, ADDRESS  
FROM    (EMPLOYEE NATURAL JOIN DEPARTMENT  
         AS DEPT(DNAME, DNO, MSSN, MSDATE))  
WHERE   DNAME='Research'
```

JOINED RELATIONS FEATURE IN SQL

- Example that illustrates multiple joins

```
SELECT    PNUMBER, DNUM, LNAME,  
FROM      (PROJECT JOIN DEPARTMENT  
           ON DNUM=DNUMBER)  
           JOIN EMPLOYEE ON MGRSSN=SSN) )  
WHERE     PLOCATION='Stafford'
```

SET OPERATIONS

- SQL has incorporated some set operations like
 - Union operation (**UNION**),
 - Set difference (**EXCEPT**) and
 - Intersection operation (**INTERSECT**)
- *Duplicate tuples are eliminated from the result*
- *Requires union compatible relations*

SET OPERATIONS (CONT.)

- Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.
- ```
(SELECT PNAME
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE DNUM=DNUMBER AND MGRSSN=SSN AND
 LNAME='Smith')

UNION

(SELECT PNAME
FROM PROJECT, WORKS_ON, EMPLOYEE
WHERE PNUMBER=PNO AND ESSN=SSN AND
 LNAME='Smith')
```

# NESTING OF QUERIES

- A complete SELECT query, called a *nested query* , can be specified within the WHERE-clause of another query, called the *outer query*
- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE
WHERE DNO IN (SELECT DNUMBER
 FROM DEPARTMENT
 WHERE DNAME='Research')
```

|   | FNAME    | LNAME   | ADDRESS                  |
|---|----------|---------|--------------------------|
| 1 | John     | Smith   | 731 Fondren, Houston, Tx |
| 2 | Franklin | Wong    | 639 Voss, Houston, TX    |
| 3 | Joyce    | English | 5631 Rice, Houston, TX   |
| 4 | Ramesh   | Narayan | 975 Fire Oak, Humble, TX |

|   | DNUMBER |
|---|---------|
| 1 | 5       |



# CORRELATED NESTED QUERIES

- If a condition in the *nested query* references an attribute of a relation declared in the *outer query* , then two queries are said to be *correlated*
- Retrieve the name of each employee who has a dependent with the same first name and gender as the employee.

```
SELECT E.FNAME, E.LNAME
FROM EMPLOYEE AS E
WHERE E.SSN IN (SELECT ESSN FROM DEPENDENT AS D
 WHERE E.SEX = D.SEX AND FNAME=DEPENDENT_NAME)
```

EMPLOYEE

| Fname    | Minit | Lname   | Ssn       | Bdate      | Address                  | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
| Alicia   | J     | Zelaya  | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
| Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
| Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
| Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
| Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
| James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M   | 55000  | NULL      | 1   |

DEPENDENT

| Essn      | Dependent_name | Sex | Bdate      | Relationship |
|-----------|----------------|-----|------------|--------------|
| 333445555 | Alice          | F   | 1986-04-05 | Daughter     |
| 333445555 | Theodore       | M   | 1983-10-25 | Son          |
| 333445555 | Joy            | F   | 1958-05-03 | Spouse       |
| 987654321 | Abner          | M   | 1942-02-28 | Spouse       |
| 123456789 | Michael        | M   | 1988-01-04 | Son          |
| 123456789 | Alice          | F   | 1988-12-30 | Daughter     |
| 123456789 | Elizabeth      | F   | 1967-05-05 | Spouse       |

**Nested query is evaluated once for each tuple in outer query**

## CORRELATED NESTED QUERIES (CONT.)

- A query written with nested SELECT... FROM... WHERE... blocks and using the = or IN comparison operators can ***always*** be expressed as a single block query.
- For example, the query on previous slide can be written as

**Retrieve the name of each employee who has a dependent with the same first name as the employee.**

|               |                                 |
|---------------|---------------------------------|
| <b>SELECT</b> | <b>E.FNAME, E.LNAME</b>         |
| <b>FROM</b>   | <b>EMPLOYEE E, DEPENDENT D</b>  |
| <b>WHERE</b>  | <b>E.SSN=D.ESSN AND</b>         |
|               | <b>E.FNAME=D.DEPENDENT_NAME</b> |

## NESTED QUERIES

---

SQL allows use of **tuples** of values in comparisons

---

- Select the Essns of all employees who work the same (project, hours) combination that 'John Smith' (whose Ssn = '123456789') works on.

- **SELECT DISTINCT** Essn

**FROM** WORKS\_ON

**WHERE** (Pno, Hours) **IN** ( **SELECT** Pno, Hours

**FROM** WORKS\_ON

**WHERE** Essn='123456789' );

# EXISTS FUNCTION

- EXISTS Function checks whether the result of a nested query is empty or not
- Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE EXISTS (SELECT *
 FROM DEPENDENT
 WHERE SSN=ESSN AND
 FNAME=DEPENDENT_NAME)
```

## EXISTS FUNCTION (CONT.)

- Retrieve the names of employees who have no dependents.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE NOT EXISTS (SELECT *
 FROM DEPENDENT
 WHERE SSN=ESSN)
```

- The above correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple.
  - If *none exist* , the EMPLOYEE tuple is selected
- EXISTS is necessary for the expressive power of SQL

## EXISTS FUNCTION (CONT.)

- Find the names of managers who have at least one dependents.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE EXISTS (SELECT *
 FROM DEPENDENT
 WHERE SSN=ESSN)

AND

EXISTS (SELECT *
 FROM DEPARTMENT
 WHERE SSN=Mgr_SSN)
```

## EXISTS FUNCTION (CONT.)

- Retrieve the name of each employee who works on *all* the projects controlled by department number 5.
- Set theory: S1 contains S2 if  $(S2 - S1 = 0)$

```

SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE NOT EXISTS (
 (SELECT PNUMBER
 FROM PROJECT
 WHERE DNUM=5)
 EXCEPT
 (SELECT PNO
 FROM WORKS_ON
 WHERE SSN=ESSN)
)

```

S1 = set of projects  
of each employee  
S2 = set of DN0=5  
projects

WORKS\_ON

| Essn      | Pno | Hours |
|-----------|-----|-------|
| 123456789 | 1   | 32.5  |
| 123456789 | 2   | 7.5   |
| 666884444 | 3   | 40.0  |
| 453453453 | 1   | 20.0  |
| 453453453 | 2   | 20.0  |
| 333445555 | 2   | 10.0  |
| 333445555 | 3   | 10.0  |
| 333445555 | 10  | 10.0  |
| 333445555 | 20  | 10.0  |
| 999887777 | 30  | 30.0  |
| 999887777 | 10  | 10.0  |
| 987987987 | 10  | 35.0  |
| 987987987 | 30  | 5.0   |
| 987654321 | 30  | 20.0  |
| 987654321 | 20  | 15.0  |
| 888665555 | 20  | NULL  |

PROJECT

| Pname           | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX        | 1       | Bellaire  | 5    |
| ProductY        | 2       | Sugarland | 5    |
| ProductZ        | 3       | Houston   | 5    |
| Computerization | 10      | Stafford  | 4    |
| Reorganization  | 20      | Houston   | 1    |
| Newbenefits     | 30      | Stafford  | 4    |

# NESTED QUERIES

CONTAINS compares two *sets* , and returns TRUE if one set contains all values in the other set.

Same as *division* operation of relational algebra

- Retrieve the name of each employee who works on *all* the projects controlled by department number 5.

- ```
SELECT  FNAME, LNAME
FROM    EMPLOYEE
WHERE   ( (SELECT  PNO
            FROM    WORKS_ON
            WHERE   SSN=ESSN)
        CONTAINS
        (SELECT  PNUMBER
            FROM    PROJECT
            WHERE   DNUM=5) )
```


NESTED CORRELATED QUERIES (CONTD)

You can also use: $s > \text{ALL } R$
 $s > \text{ANY } R$
 $\text{EXISTS } R$

Product (pname, price, category, maker)

Find products that are more expensive than all those produced
By “IBM”

```
SELECT name
FROM   Product
WHERE  price > ALL (SELECT price
                    FROM   Purchase
                    WHERE  maker='IBM')
```



NESTED CORRELATED QUERIES (CONTD)

You can also use: $s > \text{ALL } R$
 $s > \text{ANY } R$
 $\text{EXISTS } R$

Find Employee whose salary is greater than the salary of all employee in department 5

```
SELECT Fname
FROM Employee
WHERE Salary > ALL (SELECT Salary
                    FROM Employee
                    where Dno=5)
```

	Fname
1	James
2	Jennifer

COMPLEX CORRELATED QUERY

Product (pname, price, category, maker, year)

- Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

```
SELECT DISTINCT pname, maker
FROM   Product AS x
WHERE  price > ALL (SELECT price
                   FROM   Product AS y
                   WHERE  x.maker = y.maker AND x.pname!=y.pname
                   and
                   y.year < 1972)
```

Very powerful ! Also much harder to optimize.



COMPLEX CORRELATED QUERY

- Find Employee (his dno and salary) whose salary is greater than all employees in the same department

```
SELECT  Fname, Salary, Dno
FROM    Employee as E
WHERE   Salary > ALL (SELECT Salary
                      FROM    Employee as S
                      WHERE   E.dno=S.dno and E.ssn !=S.ssn )
```

	Fname	salary	Dno
1	Franklin	40000	5
2	James	55000	1
3	Jennifer	43000	4



EXPLICIT SETS

- It is also possible to use an **explicit (enumerated) set of values** in the WHERE-clause rather than a nested query
- Retrieve the social security numbers of all employees who work on project number 1, 2, or 3.

SELECT	DISTINCT ESSN
FROM	WORKS_ON
WHERE	PNO IN (1, 2, 3)

AGGREGATE FUNCTIONS

- Include **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- Find the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT      MAX(SALARY), MIN(SALARY), AVG(SALARY)
FROM        EMPLOYEE
```

- Some SQL implementations *may not allow more than one function* in the SELECT-clause

AGGREGATE FUNCTIONS (CONT.)

- Retrieve the number of employees in the 'Research' department

```
SELECT      COUNT (*)  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNO=DNUMBER AND  
            DNAME='Research'
```

GROUPING

- **GROUP BY**-clause specifies the grouping attributes
- For each department, retrieve the department number, the number of employees in the department, and their average salary.

○ **SELECT DNO, COUNT (*), AVG (SALARY)**
FROM EMPLOYEE
GROUP BY DNO

(a)

Fname	Minit	Lname	<u>Ssn</u>	...	Salary	Super_ssn	Dno
John	B	Smith	123456789		30000	333445555	5
Franklin	T	Wong	333445555		40000	9998865555	5
Ramesh	K	Narayan	666884444		38000	333445555	5
Joyce	A	English	453453453	...	25000	333445555	5
Alicia	J	Zelaya	999887777		25000	987654321	4
Jennifer	S	Wallace	987654321		43000	9998865555	4
Ahmad	V	Jabbar	987987987		25000	987654321	4
James	E	Bong	9998865555		55000	NULL	1

Dno	Count (*)	Avg (Salary)
5	4	33250
4	3	31000
1	1	55000

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

GROUPING (CONT.)

- For each project, retrieve the project number, project name, and the number of employees who work on that project.



```
SELECT PNUMBER, PNAME, COUNT (*)  
FROM PROJECT, WORKS_ON  
WHERE PNUMBER=PNO  
GROUP BY PNUMBER, PNAME
```

	PNUMBER	PNAME	count
1	1	ProductX	2
2	2	ProductY	3
3	3	ProductZ	2
4	10	Compu...	3
5	20	Reorga...	3
6	30	Newbe...	3

- The grouping and functions are applied *after* the joining of the two relations
- Group By clause specifies grouping attributes which should appear in SELECT clause

HAVING-CLAUSE

- HAVING-clause specify a selection condition on groups
- For each project *on which more than two employees work* , retrieve the project number, project name, and the number of employees who work on that project.



SELECT	PNUMBER, PNAME, COUNT (*)
FROM	PROJECT, WORKS_ON
WHERE	PNUMBER=PNO
GROUP BY	PNUMBER, PNAME
HAVING	COUNT (*) > 2

HAVING-CLAUSE

(b)

Pname	Pnumber	...	Essn	Pno	Hours
ProductX	1		123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10	...	333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

PNUMBER	PNAME	count
2	ProductY	3
10	Computerization	3
20	Reorganization	3
30	Newbenefits	3

After applying the WHERE clause but before applying HAVING

HAVING-CLAUSE

----- Applying the HAVING clause condition -----

Pname	Pnumber	...	Easn	Pno	Hours		Pname	Count (*)
ProductY	2		123456789	2	7.5	}	ProductY	3
ProductY	2		453453453	2	20.0		Computerization	3
ProductY	2		333445555	2	10.0		Reorganization	3
Computerization	10		333445555	10	10.0	}	Newbenefits	3
Computerization	10	...	999887777	10	10.0			
Computerization	10		987987987	10	35.0	}		
Reorganization	20		333445555	20	10.0			
Reorganization	20		987654321	20	15.0			
Reorganization	20		888665555	20	NULL	}		
Newbenefits	30		987987987	30	5.0			
Newbenefits	30		987654321	30	20.0			
Newbenefits	30		999887777	30	30.0			

Result of Q26
(Pnumber not shown)

After applying the HAVING clause condition


GROUP BY AND HAVING

- Count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work
- **SELECT** Dname, **COUNT** (*)
FROM DEPARTMENT, EMPLOYEE
WHERE Dnumber=Dno **AND** Salary>40000
GROUP BY Dname
HAVING COUNT (*) > 5;

GENERAL FORM OF GROUPING AND AGGREGATION

```
SELECT  S
FROM    R1,...,Rn
WHERE   C1
GROUP BY a1,...,ak
HAVING  C2
```

Evaluation steps:

1. Evaluate FROM-WHERE, apply condition C1
 2. Group by the attributes a_1, \dots, a_k
 3. Apply condition C2 to each group (may have aggregates)
 4. Compute aggregates in S and return the result
- 

TWO EXAMPLES

Store(sid, sname)

Product(pid, pname, price, sid)

Find stores that sell *only* products with price > 100

same as:

Find stores s.t. all their products have price > 100)



```
SELECT Store.name
FROM Store, Product
WHERE Store.sid = Product.sid
GROUP BY Store.sid, Store.name
HAVING 100 < min(Product.price)
```

Find stores s.t. all their
products have price > 100

Almost equivalent...

```
SELECT Store.name
FROM Store
WHERE
    100 < ALL (SELECT Product.price
               FROM product
               WHERE Store.sid = Product.sid)
```

```
SELECT Store.name
FROM Store
WHERE Store.sid NOT IN
    (SELECT Product.sid
     FROM Product
     WHERE Product.price <= 100)
```



AGGREGATE EXAMPLE

- **Example:** Count the number of distinct salary values in the database.
 - **SELECT COUNT (DISTINCT Salary)**
 - **FROM EMPLOYEE;**
- NULL values are **discarded** when aggregate functions are applied to a particular attribute.

SUMMARY OF SQL QUERIES

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order:

SELECT <attribute list>
FROM <table list>
[WHERE <condition>]
[GROUP BY <grouping attribute(s)>]
[HAVING <group condition>]
[ORDER BY <attribute list>]

- A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause

SUMMARY OF SQL QUERIES (CONT.)

- The SELECT-clause lists the attributes or functions to be retrieved
- The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The WHERE-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- GROUP BY specifies grouping attributes
- HAVING specifies a condition for selection of groups
- ORDER BY specifies an order for displaying the result of a query

SQL QUERIES

- There are various ways to specify the same query in SQL
 - This is to give flexibility to user to specify queries
- For query optimization, it is preferable to write a query with as little nesting and implied ordering as possible.
- Ideally, DBMS should process the same query in the same way regardless of how the query is specified.
 - But this is quite difficult in practice, (chapter 19,20)

SPECIFYING UPDATES IN SQL

- There are three SQL commands to modify the database;
 - INSERT,
 - DELETE, and
 - UPDATE

INSERT

- It is used to add one or more tuples to a relation

- **Example:**

```
INSERT INTO EMPLOYEE  
VALUES ('Richard','K','Marini', '653298653', '30-DEC-52',  
'98 Oak Forest,Katy,TX', 'M', 37000,'987654321', 4 )
```

- Attribute values should be listed in the same order as the attributes were specified in the CREATE TABLE command

INSERT (CONT.)

- An alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple
- **Example:** Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

```
INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)  
VALUES ('Richard', 'Marini', '653298653')
```

- Attributes with NULL values can be left out

INSERT (CONT.)

- Suppose we want to create a temporary table that has the name, number of employees, and total salaries for each department.
- A table DEPTS_INFO is created by Q1, and is loaded with the information retrieved from the database by the query Q2.

Q1: **CREATE TABLE DEPTS_INFO**
 (DEPT_NAME VARCHAR(10),
 NO_OF_EMPS INTEGER,
 TOTAL_SAL INTEGER);

Q2: **INSERT INTO DEPTS_INFO (DEPT_NAME,**
 NO_OF_EMPS, TOTAL_SAL)
 SELECT **DNAME, COUNT (*), SUM (SALARY)**
 FROM **DEPARTMENT, EMPLOYEE**
 WHERE **DNUMBER=DNO**
 GROUP BY **DNAME ;**

DELETE

- Removes tuples from a relation
- Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
- Examples:

```
DELETE FROM EMPLOYEE
WHERE LNAME='Brown'
```

```
DELETE FROM EMPLOYEE
WHERE DNO IN (SELECT DNUMBER
FROM DEPARTMENT
WHERE DNAME='Research')
```

```
DELETE FROM EMPLOYEE
```

UPDATE

- Used to modify attribute values of selected tuples
- **Example:** Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
UPDATE PROJECT  
SET      PLOCATION = 'Bellaire', DNUM = 5  
WHERE   PNUMBER=10
```

UPDATE (CONT.)

- **Example:** Give all employees in the 'Research' department a 10% raise in salary.

```
UPDATE EMPLOYEE
SET      SALARY = SALARY * 1.1
WHERE DNO IN (SELECT DNUMBER
                FROM DEPARTMENT
                WHERE DNAME='Research')
```

VIEWS IN SQL

- A view is a “virtual” table that is derived from other tables
- Allows for limited update operations (since the table may not physically be stored)
- Allows full query operations
- A convenience for expressing certain operations
- They are used to:
 - simplify complex queries, and
 - define distinct conceptual interfaces for different users.



SQL VIEWS: AN EXAMPLE

- Specify a different WORKS_ON table

CREATE VIEW WORKS_ON1 AS

SELECT FNAME, LNAME, PNAME, HOURS
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN=ESSN AND PNO=PNUMBER

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------



SQL VIEWS: AN EXAMPLE2

```
CREATE VIEW    DEPT_INFO(Dept_name, No_of_emps, Total_sal)
AS SELECT     Dname, COUNT (*), SUM (Salary)
FROM          DEPARTMENT, EMPLOYEE
WHERE         Dnumber=Dno
GROUP BY     Dname;
```

DEPT_INFO

Dept_name	No_of_emps	Total_sal
-----------	------------	-----------

USING A VIRTUAL TABLE

- We can specify SQL queries on a newly created view:

```
SELECT FNAME, LNAME  
FROM WORKS_ON1  
WHERE PNAME='ProductX';
```

- DBMS is responsible to keep view always up-to-date
- When no longer needed, a view can be dropped:

```
DROP WORKS_ON1;
```



EFFICIENT VIEW IMPLEMENTATION

- **Query modification:** present the view query in terms of a query on the underlying base tables

```
SELECT    Fname, Lname  
FROM      WORKS_ON1  
WHERE     Pname='ProductX';
```

```
SELECT FNAME, LNAME, PNAME, HOURS  
FROM EMPLOYEE, PROJECT, WORKS_ON  
WHERE SSN=ESSN AND PNO=PNUMBER AND PNAME='PRODUCTX'
```

- Disadvantage:
 - Inefficient for views defined via complex queries (esp if additional queries are to be applied within a short time period)



EFFICIENT VIEW IMPLEMENTATION

- **View materialization:** involves physically creating and keeping a temporary table
 - assumption: other queries on the view will follow
 - concerns: maintaining correspondence between the base table and the view when the base table is updated
 - strategy: incremental update



VIEW UPDATE

- **Single view without aggregate operations:**

- update may map to an update on the underlying base table

- **Views involving joins:**

- an update *may* map to an update on the underlying base relations
- not always possible

- **Example:**

UPDATE WORKS_ON1

SET PNAME=XYZ'

WHERE FNAME='JOHN AND
LNAME='SMITH' AND
PNAME='PRODUCTX'

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4



A) UPDATE WORKS_ON1

SET PNO = (SELECT PNUMBER FROM PROJECT
WHERE PNAME='XYZ')

WHERE ESSN IN (SELECT SSN FROM EMPLOYEE
WHERE LNAME='SMITH' AND FNAME='JOHN')

AND

PNO = (SELECT PNUMBER FROM PROJECT
WHERE PNAME='PRODUCTX')

B) UPDATE PROJECT SET PNAME='XYZ'

WHERE PNAME='PRODUCTX'

UN-UPDATABLE VIEWS

- Views defined using groups and aggregate functions are not updateable

```
UPDATE DEPT_INFO  
SET      Total_sal=100000  
WHERE    Dname='Research';
```

- Views defined on multiple tables using joins are generally not updateable
- WITH CHECK OPTION: must be added to the definition of a view if the view is to be updated
 - to allow check for updatability and to plan for an execution strategy



ASSERTION

- In SQL, table constraints are associated with a single table.
 - **CHECK(DNUMBER >0 AND DNUMBER <25)**
 - **CHECK (Dept_create_date <= Mgr_start_date)**
- Expression in the **CHECK** clause can refer to other tables but when a constraint involves many tables, it becomes cumbersome
- Assertions are constraints that are not associated with any one table.
- SQL-server do not support assertion, however it do support trigger and check constraint

ASSERTIONS: AN EXAMPLE

- “The salary of an employee must not be greater than the salary of the manager of the department that the employee works for”

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK (NOT EXISTS (SELECT *
                    FROM EMPLOYEE E, EMPLOYEE M, DEPARTMENT D
                    WHERE E.SALARY > M.SALARY AND
                          E.DNO=D.NUMBER AND D.MGRSSN=M.SSN))
```



USING GENERAL ASSERTIONS

- Specify a query that violates the condition; include inside a `NOT EXISTS` clause
- Query result must be empty
 - if the query result is not empty, the assertion has been violated
- **Create Assertion to enforce the** constraint that
 - The number of Employees in each department should be less than 20.
 - An employee should be assign a project of his\her department only.
 - The total working hours of all the employees in a department should be less than 180.



- CHECK clause can also be used to specify constraints on *individual* attributes, domains and *individual* tuples

CONSTRAINT deptPK PRIMARY KEY(Dnumber)

CONSTRAINT deptSK UNIQUE(Dname)

CHECK (Dept_create_date <= Mgr_start_date)

- Difference between CREATE ASSERTION and the individual domain and tuple constraints is
 - CHECK on individual attributes, domains, and tuples are checked *only when tuples are inserted or updated*.

SQL TRIGGERS

- Objective: to monitor a database and take action when a condition occurs
- Triggers are expressed in a syntax similar to assertions and include the following:
 - event (e.g., an update operation)
 - condition
 - action (to be taken when the condition is satisfied)



SQL TRIGGERS: AN EXAMPLE

- A trigger to compare an employee's salary to his/her supervisor during insert or update operations:

```
CREATE TRIGGER INFORM_SUPERVISOR
BEFORE INSERT OR UPDATE OF
    SALARY, SUPERVISOR_SSN ON EMPLOYEE
FOR EACH ROW
    WHEN
        (NEW.SALARY > (SELECT SALARY FROM EMPLOYEE
                        WHERE SSN=NEW.SUPERVISOR_SSN))
INFORM_SUPERVISOR (NEW.SUPERVISOR_SSN, NEW.SSN;
```



EXAMPLE: BOAT RENTAL DATABASE

- Consider the following **Boat Rental database** schema
 - **SAILOR** (SID, SName, Phone, City)
 - **BOAT** (BName, BType, Price, OID)
 - **RESERVATION** (SID, BName, Date, Duration)
 - **OWNER** (OID, OName, Phone, Street, City, Country)
- Find the name and city of sailors who reserve every 'Fishing' boat.
- Find the names of boats that are reserved by at least ten different sailors.
- List names, owner names, and prices of the boats which were reserved in 2007 but not in 2008.

EXAMPLE: BOAT RENTAL DATABASE

- Consider the following schema
 - SAILOR (SID, SName, Phone, City)
 - BOAT (BName, BType, Price, OID)
 - RESERVATION (SID, BName, Date, Duration)
 - OWNER (OID, OName, Phone, Street, City, Country)
- Find the name and city of sailors who reserve every 'Fishing' boat.

```
○ SELECT SNAME, CITY
  FROM SAILOR s
 WHERE NOT EXISTS ( (SELECT BNAME
                     FROM BOAT
                     WHERE BTYPE='Fishing')
                   EXCEPT
                   (SELECT BNAME
                    FROM reservation r
                    WHERE r.SID=s.SID) )
```

EXAMPLE: BOAT RENTAL DATABASE

- Consider the following schema
 - SAILOR (SID, SName, Phone, City)
 - BOAT (BName, BType, Price, OID)
 - RESERVATION (SID, BName, Date, Duration)
 - OWNER (OID, OName, Phone, Street, City, Country)
- **Find the names of boats that are reserved by at least ten different sailors.**
- - Select bname
 - From reservation r
 - Group by bname
 - Having count(DISTINCT SID) >9

EXAMPLE: BOAT RENTAL DATABASE

- Consider the following schema
 - SAILOR (SID, SName, Phone, City)
 - BOAT (BName, BType, Price, OID)
 - RESERVATION (SID, BName, Date, Duration)
 - OWNER (OID, OName, Phone, Street, City, Country)
- **List names, owner names, and prices of the boats which were reserved in 2007 but not in 2008.**
- **Select** distinct b.bname, b.price, o.ename
 - From** reservation r, boat b, owner o
 - Where** r.bname = b.bname and b.oid=o.oid and r.date LIKE '%2007%' and r.sid not in
 - (Select ic.sid
 - From reservation ic, boat ip
 - Where ic.bname = ip.bname and ip.date LIKE '%2008%')

EXAMPLE: BOAT RENTAL DATABASE

- Consider the following Boat Rental database schema:

- SAILOR (SID, SName, Phone, City)
- BOAT (BName, BType, Price, OID)
- RESERVATION (SID, BName, Date, Duration)
- OWNER (OID, OName, Phone, Street, City, Country)

- SELECT DISTINCT Bname

- FROM BOAT

- WHERE Price > ALL (SELECT price
FROM BOAT b , OWNER o
WHERE b.oid=o.oid and
Country='Pakistan')

What does the query do?



EXAMPLE: BOAT RENTAL DATABASE

- **Consider the following Boat Rental database schema:**

- SAILOR (SID, SName, Phone, City)
- BOAT (BName, BType, Price, OID)
- RESERVATION (SID, BName, Date, Duration)
- OWNER (OID, OName, Phone, Street, City, Country)

- Select bname,count(*)
- From reservation r ,boat b,owner o
- Where b.bname=r.bname and b. oid=o.oid and country='Pakistan'
- Group by bname
- Having count(*) > 5

What does the query do?

